

# NLP Corpus Analysis Tools

## Project Overview

This project provides two Python scripts for analyzing text corpora and performing various natural language processing (NLP) tasks. The scripts apply n-gram models to analyze word frequencies, calculate probabilities, implement smoothing techniques, and evaluate language models through perplexity.

## Features

- Unigram and bigram frequency analysis
- Maximum Likelihood Estimation (MLE) for probability calculation
- Add-one (Laplace) smoothing implementation
- Perplexity calculation for language model evaluation
- Random sentence generation based on n-gram models
- Interactive menu-based interface

## Requirements

- Python 3.10 or newer
- No external dependencies (uses only standard library modules):
  - argparse for command-line argument parsing
  - collections for efficient counting
  - random for sampling operations
  - math for logarithmic calculations

## File Structure

project\_root/

```
|— ForAnalysis.py    # Script with additional option to print top 10 n-grams
|— Updated_T1-T4.py  # Base script for corpus analysis and NLP tasks
|— TheStory.txt      # Text corpus for analysis
|— README.txt        # This documentation file
```

# Installation

Simply download the files and ensure Python 3.10+ is installed on your system.

## Usage Instructions

### Basic Command Syntax

```
python3 Updated_T1-T4.py corpus_path
```

```
python3 ForAnalysis.py corpus_path
```

Where `corpus_path` is the path to your text corpus file (e.g., `TheStory.txt`).

### Interactive Options

After running either script, an interactive menu will be displayed with the following options:

#### For Updated\_T1-T4.py:

- 1: Reprint probabilities
- 2: Generate sentence
- 3: Apply Add-1 smoothing
- 4: Calculate Perplexity on test data
- 5: Quit

#### For ForAnalysis.py:

- 1: Reprint probabilities
- 2: Generate sentence
- 3: Apply Add-1 smoothing
- 4: Calculate Perplexity on test data
- 5: Quit

6: Print top 10 uni|bi grams

## Example Usage

1. Run the basic analysis script:  
`python3 Updated_T1-T4.py TheStory.txt`
2. Run the extended analysis script (includes top 10 n-grams option):  
`python3 ForAnalysis.py TheStory.txt`
3. For perplexity calculation (Option 4), you'll be prompted to enter the path to a test corpus file.

## Menu Options Explained

### Option 1: Reprint probabilities

Displays the first five unigram and bigram probabilities calculated from the corpus.

### Option 2: Generate sentence

Creates a random 15-word sentence using the calculated n-gram model. If smoothing has been applied, the sentence will use the smoothed probabilities.

### Option 3: Apply Add-1 smoothing

Applies Laplace (add-one) smoothing to the probability distributions to handle zero-probability n-grams.

### Option 4: Calculate Perplexity on test data

Prompts for a test corpus path and calculates the perplexity score, which evaluates how well the n-gram model predicts the test data.

### Option 5: Quit

Exits the program.

### Option 6: Print top 10 uni|bi grams (ForAnalysis.py only)

Displays the top 10 most probable unigrams and bigrams in the corpus.

## Technical Details

## Tokenization

The scripts use a custom tokenizer that identifies words as sequences of alphanumeric characters and underscores.

## Probability Calculation

- Unigram probabilities are calculated as:  $P(w) = \text{count}(w) / N$  (where  $N$  is the total number of tokens)
- Bigram probabilities are calculated as:  $P(w_2|w_1) = \text{count}(w_1, w_2) / \text{count}(w_1)$

## Add-One Smoothing

Implements Laplace smoothing where:

- Smoothed unigram probability:  $P(w) = (\text{count}(w) + 1) / (N + V)$
- Smoothed bigram probability:  $P(w_2|w_1) = (\text{count}(w_1, w_2) + 1) / (\text{count}(w_1) + V)$  Where  $V$  is the vocabulary size.

## Perplexity

Calculated as  $2^{(-L)}$ , where  $L$  is the average log (base 2) probability of the test corpus according to the language model.

## Sentence Generation

Uses the calculated probability distributions to randomly generate sentences, with words selected proportionally to their likelihood in the model.